

1. Section 4.7, page 168, question 1

2. Section 4.7, page 169-170, question 5

a.

If the Bayes decision boundary is linear, we expect LDA to perform better on the testing set because it is much less flexible and performs well in situations where having lower variance is important, like if Bayes decision boundary is linear. In contrast, QDA performs poorly well on the training set because it models the random error of the data better, but this will make it perform worse on the test data.

b.

If the Bayes decision boundary is non-linear, we expect QDA to perform better on the training and test sets because it has higher variance, which would approximate non-linearities in the data better. In contrast, LDA, as it is a linear method, would result in high bias across the data.

c.

As the sample size increases, we expect the test prediction accuracy of both QDA and LDA to increase. It depends which one will perform better—if the data is non-linear, then QDA should perform better. If the data is linear, then LDA will perform better. Also, if the sample size doesn't increase too much (stays relatively small), then LDA could perform better than QDA because QDA needs more samples to perform well.

d.

True, even if the Bayes decision boundary for a given problem is linear, we will probably achieve a superior test error rate using QDA rather than LDA because QDA is flexible enough to model a linear decision boundary.

While the statement is true, you have to worry about overfitting. QDA will not only be able to model the linear decision boundary, but it can also capture some of the residual error. This will help it perform well on the training data and possibly the test data since it is more flexible than LDA. If QDA models too much of the residual error in the training data (overfit), it might perform much worse on the test data.

3. Section 4.7, page 170, question 8

Logistic regression: training error = 20%, testing error = 30%

KNN (k=1): training error = ?, testing error = ?, average error = 18%

KNN with $k=1$ will classify all the data in the training set correctly (training error = 0%), but the test set will almost certainly have errors. With the knowledge that the training error for KNN is equal to zero, we can use a sample dataset to compare the testing error of the logistic regression and KNN. Assume there is a dataset with 1,000 samples so we will have 500 training and 500 testing samples. The overall error for KNN is 18%, which means there are 180 misclassified samples for KNN. Since the training error of KNN is equal to zero, then these 180 samples had to come from testing. We can calculate KNN's testing error to be 36% (180 misclassified samples / 500 testing samples). As a result, you can clearly see that logistic regression's testing error (30%) is better than KNN's (36%) so you would prefer logistic regression for classifying new observations in the problem.

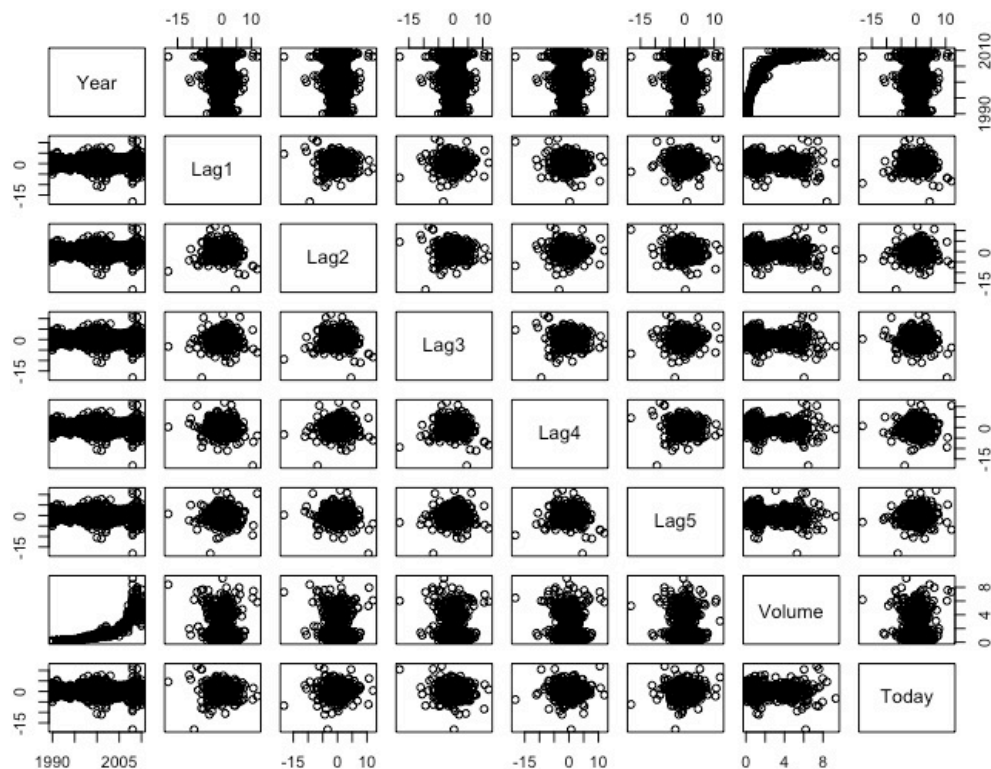
4. Section 4.7, page171, question 10
a.

```
> summary(Weekly)
```

Year	Lag1	Lag2	Lag3	Lag4
Min. :1990	Min. :-18.1950	Min. :-18.1950	Min. :-18.1950	Min. :-18.1950
1st Qu.:1995	1st Qu.: -1.1540	1st Qu.: -1.1540	1st Qu.: -1.1580	1st Qu.: -1.1580
Median :2000	Median : 0.2410	Median : 0.2410	Median : 0.2410	Median : 0.2380
Mean :2000	Mean : 0.1506	Mean : 0.1511	Mean : 0.1472	Mean : 0.1458
3rd Qu.:2005	3rd Qu.: 1.4050	3rd Qu.: 1.4090	3rd Qu.: 1.4090	3rd Qu.: 1.4090
Max. :2010	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260


Lag5	Volume	Today	Direction
Min. :-18.1950	Min. :0.08747	Min. :-18.1950	Down:484
1st Qu.: -1.1660	1st Qu.:0.33202	1st Qu.: -1.1540	Up :605
Median : 0.2340	Median :1.00268	Median : 0.2410	
Mean : 0.1399	Mean :1.57462	Mean : 0.1499	
3rd Qu.: 1.4050	3rd Qu.:2.05373	3rd Qu.: 1.4050	
Max. : 12.0260	Max. :9.32821	Max. : 12.0260	

Based on an initial numerical summary of the data, there does not appear to be any pattern. All the lags (Lag1-Lag5) appear very consistent with each other. Further investigating, I looked at a scatterplot matrix of all the quantitative variables.



Unfortunately, it is very difficult to read—not only because there is a lot of data, but also because there aren't very main patterns to pick out. One pattern that looks promising though is the relationship between year and volume.

Exploring the relationship further, it appears that over the years, the trade volume has increased visually significantly.



Looking into the correlation matrix for the dataset, clearly confirms that there is a significant, positive relationship between year and volume, whereas all the other correlations are very small.

```
> cor(weekly_quantitative)
```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today
Year	1.00000000	-0.032289274	-0.03339001	-0.03000649	-0.031127923	-0.030519101	0.84194162	-0.032459894
Lag1	-0.03228927	1.000000000	-0.07485305	0.05863568	-0.071273876	-0.008183096	-0.06495131	-0.075031842
Lag2	-0.03339001	-0.074853051	1.000000000	-0.07572091	0.058381535	-0.072499482	-0.08551314	0.059166717
Lag3	-0.03000649	0.058635682	-0.07572091	1.000000000	-0.075395865	0.060657175	-0.06928771	-0.071243639
Lag4	-0.03112792	-0.071273876	0.05838153	-0.07539587	1.000000000	-0.075675027	-0.06107462	-0.007825873
Lag5	-0.03051910	-0.008183096	-0.07249948	0.06065717	-0.075675027	1.000000000	-0.05851741	0.011012698
Volume	0.84194162	-0.064951313	-0.08551314	-0.06928771	-0.061074617	-0.058517414	1.000000000	-0.033077783
Today	-0.03245989	-0.075031842	0.05916672	-0.07124364	-0.007825873	0.011012698	-0.03307778	1.000000000

b.

Yes, Lag1 appears to be statistically significant at the 0.05 level. The intercept of the model is also significant.

```

> glm.fit=glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume, data=Weekly, family=binomial)
> summary(glm.fit)

Call:
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
    Volume, family = binomial, data = Weekly)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.6949  -1.2565   0.9913   1.0849   1.4579

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept)  0.26686    0.08593   3.106  0.0019 **
Lag1        -0.04127    0.02641  -1.563  0.1181
Lag2         0.05844    0.02686   2.175  0.0296 *
Lag3        -0.01606    0.02666  -0.602  0.5469
Lag4        -0.02779    0.02646  -1.050  0.2937
Lag5        -0.01447    0.02638  -0.549  0.5833
Volume      -0.02274    0.03690  -0.616  0.5377
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 1496.2  on 1088  degrees of freedom
Residual deviance: 1486.4  on 1082  degrees of freedom
AIC: 1500.4

Number of Fisher Scoring iterations: 4

```

c.

```

glm.probs = predict(glm.fit, Weekly, type="response")
glm.pred = rep("Down", nrow(Weekly))
glm.pred[glm.probs > 0.50] = "Up"
table(glm.pred, Direction)
mean(glm.pred == Direction)

```

	Direction	
glm.pred	Down	Up
Down	54	48
Up	430	557

Based on the confusion matrix, the model does not fit the data very well. The model predicts correctly 54 downs and 557 ups in the market, but incorrectly predicts 430 downs and 48 ups. Overall, this means that the model predicts the market direction only 56.1% of the time on the training set. Training set predictions tend to be overly optimistic—when using a test set, performance will likely decrease.

d.

```

train = Year <= 2008
Weekly2 = Weekly[!train,]
glm.fit = glm(Direction~Lag2, family=binomial, data=Weekly, subset=train)
summary(glm.fit)
glm.probs = predict(glm.fit, Weekly2, type="response")
glm.pred = rep("Down", nrow(Weekly2))
glm.pred[glm.probs > 0.50] = "Up"
table(glm.pred, Weekly2$Direction)
mean(glm.pred == Weekly2$Direction)

```

	Down	Up
glm.pred	Down	Up
Down	9	5
Up	34	56

Overall fraction of correct predictions = 0.625

This model appears to fit the data better, even though we used two completely separate training and testing datasets.

e.

```
library(MASS)
lda.fit = lda(Direction~Lag2, data=Weekly, subset=train)
lda.predict = predict(lda.fit, Weekly2)
lda.class = lda.predict$class
table(lda.class, Weekly2$Direction)
mean(lda.class == Weekly2$Direction)
```

lda.class	Down	Up
Down	9	5
Up	34	56

Overall fraction of correct predictions = 0.625

LDA performs identical to the previous logistic regression

f.

```
qda.fit = qda(Direction~Lag2, data=Weekly, subset=train)
qda.fit
qda.predict = predict(qda.fit, Weekly2)
qda.class = qda.predict$class
table(qda.class, Weekly2$Direction)
mean(qda.class == Weekly2$Direction)
```

qda.class	Down	Up
Down	0	0
Up	43	61

Overall fraction of correct predictions = 0.586

QDA performs poorly on the test data. It predicts all the down directions incorrectly, but it predicts all the up directions correctly. The model could be overfitting the up direction and is unable to approximate anything else.

g.

```
library(class)
train.X = Weekly[train, "Lag2", drop=F]
test.X = Weekly[!train, "Lag2", drop=F]
train.Direction = Weekly[train, "Direction", drop=T]
test.Direction = Weekly[!train, "Direction", drop=T]
set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k=1)
table(knn.pred, test.Direction)
mean(knn.pred == test.Direction)
```

	test.Direction	
knn.pred	Down	Up
Down	21	30
Up	22	31

Overall fraction of correct predictions = 0.5

KNN does not appear very effective at modeling the data at all.

h.

The Logistic Regression (d) and Linear Discriminate Analysis (e) appear to perform the best on the data. Quadratic Discriminant Analysis (f) appears to overfit the up direction in the data, while K-Nearest Neighbors looks completely random, almost like guessing could have performed better.

i.

Logistic regression with predictors: Lag2, Volume, and Lag2*Volume

```
glm.fit = glm(Direction~Lag2*Volume, family=binomial, data=Weekly, subset=train)
summary(glm.fit)
glm.probs = predict(glm.fit, Weekly2, type="response")
glm.pred = rep("Down", nrow(Weekly2))
glm.pred[glm.probs > 0.50] = "Up"
table(glm.pred, Weekly2$Direction)
mean(glm.pred == Weekly2$Direction)
```

glm.pred	Down	Up
Down	20	25
Up	23	36

Overall fraction of correct predictions = 0.538

Does not perform very well.

KNN with predictor: Lag2, and k=4

```
set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k=4)
table(knn.pred, test.Direction)
mean(knn.pred == test.Direction)
```

	test.Direction	
knn.pred	Down	Up
Down	20	17
Up	23	44

Overall fraction of correct predictions = 0.615

Performs much better than KNN with k=1, but does not perform as well as the original Logistic Regression or Linear Discriminant Analysis.

KNN with predictor: Lag2, Lag2^2, and k=4

```
train.X = cbind(Lag2, I(Lag2 ^ 2))[train,]
test.X = cbind(Lag2, I(Lag2 ^ 2))[!train,]
train.Direction = Direction [train]
set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k=4)
table(knn.pred, test.Direction)
mean(knn.pred == test.Direction)
```

	test.Direction	
knn.pred	Down	Up
Down	20	17
Up	23	44

Overall fraction of correct predictions = 0.615

Performs the same as the previous KNN model, even with a non-linear transformation on Lag2.

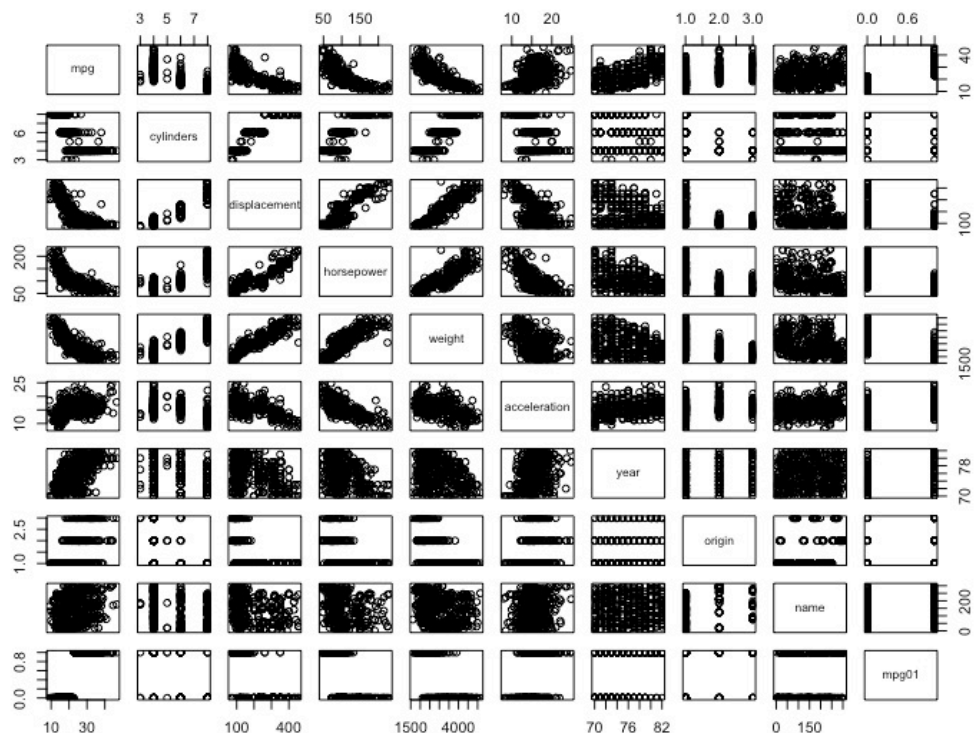
5. Section 4.7, page 171-172, question 11

a.

```
library(ISLR)
data(Auto)
Auto$mpg01 <- ifelse(Auto$mpg > median(Auto$mpg), 1, 0)
```

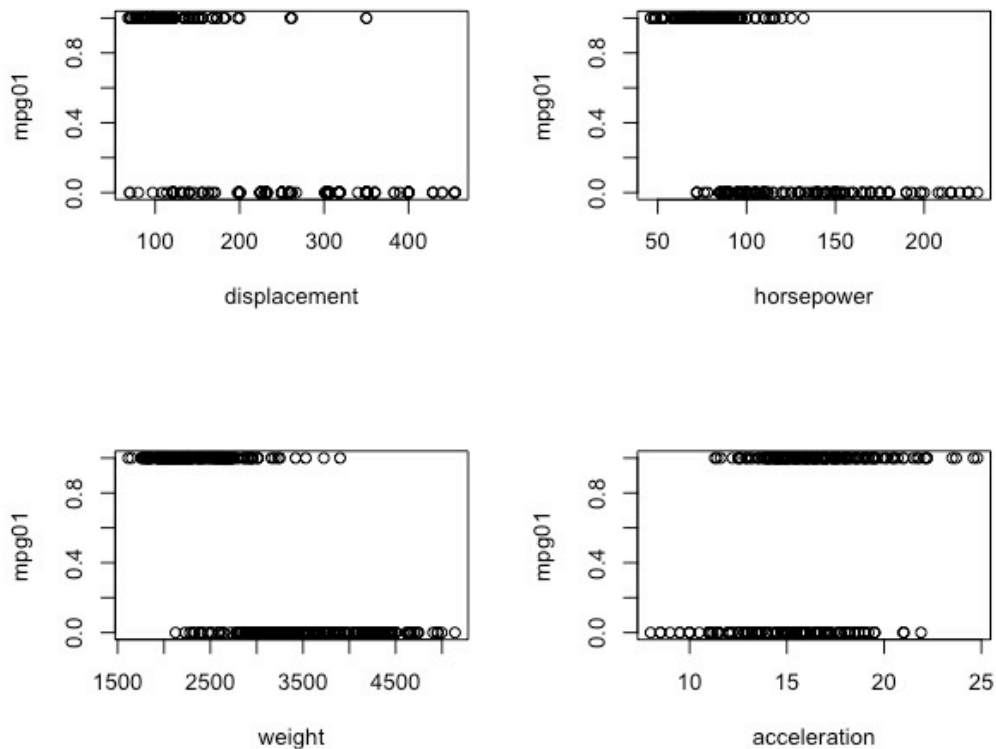
b.

```
pairs(Auto)
```



Displacement, horsepower, weight, and acceleration all seem to be useful in predicting mpg01 because there is a pattern to the data that looks like it can be modeled through classification. Investigating further below, we can see the relationships better than in using a scatterplot matrix.

```
par(mfrow=c(2, 2))
plot(displacement, mpg01)
plot(horsepower, mpg01)
plot(weight, mpg01)
plot(acceleration, mpg01)
```



c.

```
set.seed(1)
rands <- rnorm(nrow(Auto))
test <- rands > quantile(rands, 0.75)
train <- !test
Auto.train <- Auto[train,]
Auto.test <- Auto[test,]
```

d.

```
library(MASS)
lda.fit = lda(mpg01~displacement+horsepower+weight+acceleration, data=Auto,
subset=train)
lda.fit
lda.predict = predict(lda.fit, Auto.test)
lda.class = lda.predict$class
table(lda.class, Auto.test$mpg01)
test_accuracy <- mean(lda.class == Auto.test$mpg01)
1 - test_accuracy
```

```
lda.class  0  1
          0 33  3
          1 11 51
```

The test accuracy of LDA is 0.8571429, while the test error is 0.1428571.

e.

```
qda.fit = qda(mpg01~displacement+horsepower+weight+acceleration, data=Auto.train)
qda.fit
qda.predict = predict(qda.fit, Auto.test)
qda.class = qda.predict$class
table(qda.class, Auto.test$mpg01)
test_accuracy <- mean(qda.class == Auto.test$mpg01)
1 - test_accuracy
```

```
qda.class  0  1
          0 36  5
          1  8 49
```

The test accuracy of QDA is 0.8673469, while the test error is 0.1326531.

f.

```
glm.fit = glm(mpg01~displacement+horsepower+weight+acceleration, family=binomial,
data=Auto.train)
summary(glm.fit)
glm.probs = predict(glm.fit, Auto.test, type="response")
glm.pred = rep(0, nrow(Auto.test))
glm.pred[glm.probs > 0.50] = 1
table(glm.pred, Auto.test$mpg01)
test_accuracy <- mean(glm.pred == Auto.test$mpg01)
1 - test_accuracy
```

```
glm.pred  0  1
          0 36  6
          1  8 48
```

The test accuracy of logistic regression is 0.8571429, while the test error is 0.1428571.

g.

```
library(class)
set.seed(1)
train.knn = Auto.train[,c("displacement", "horsepower", "weight", "acceleration")]
test.knn = Auto.test[,c("displacement", "horsepower", "weight", "acceleration")]
```

k=1

```
knn.pred=knn(train.knn, test.knn, Auto.train$mpg01, k=1)
table(knn.pred, Auto.test$mpg01)
test_accuracy <- mean(knn.pred==Auto.test$mpg01)
1 - test_accuracy
```

```
knn.pred  0  1
          0 35  5
          1  9 49
```

The test accuracy of KNN (k=1) is 0.8571429, while the test error is 0.1428571.

k=2

```
knn.pred=knn(train.knn, test.knn, Auto.train$mpg01, k=2)
table(knn.pred, Auto.test$mpg01)
test_accuracy <- mean(knn.pred==Auto.test$mpg01)
1 - test_accuracy
```

```
knn.pred  0  1
          0 34  5
          1 10 49
```

The test accuracy of KNN (k=2) is 0.8469388, while the test error is 0.1530612.

k=3 performs the best on the data set

```
knn.pred=knn(train.knn, test.knn, Auto.train$mpg01, k=3)
table(knn.pred, Auto.test$mpg01)
test_accuracy <- mean(knn.pred==Auto.test$mpg01)
1 - test_accuracy
```

```
knn.pred  0  1
          0 36  4
          1  8 50
```

The test accuracy of KNN (k=3) is 0.877551, while the test error is 0.122449.

k=4

```
knn.pred=knn(train.knn, test.knn, Auto.train$mpg01, k=4)
table(knn.pred, Auto.test$mpg01)
test_accuracy <- mean(knn.pred==Auto.test$mpg01)
1 - test_accuracy
```

```
knn.pred  0  1
          0 36  5
          1  8 49
```

The test accuracy of KNN (k=4) is 0.8673469, while the test error is 0.1326531.

6. Section 5.4, page 197, question 1

7. Section 5.4, page 197, question 2

a.

Bootstrap is sampling with replacement so even when we start with n samples, the sample size will be the same every time because the observation is returned to the sample n . The first bootstrap observation, j , is one of the n samples so there are $n - 1$ samples that are not j . Thus, we get the following:

b.

The second bootstrap observation has the same probability that it is not the j th observation from the original sample as **a** because we sampling with replacement!

c.

d.

```
prob_in_boot.fn = function (n) {  
  return (1 - ((1 - 1/n)^n))  
}
```

```
prob_in_boot.fn(5)
```

0.67232

e.

```
prob_in_boot.fn(100)
```

0.6339677


f.

```
prob_in_boot.fn(10000)
```

0.632139

g.

```
x = seq(1, 100000)
y = sapply(x, function (n) { prob_in_boot.fn(n) })
plot(x, y, xlab="n", ylab="probability", log="x")
```



As n increases, the probability decreases, until n is around 100. After that the probability stays pretty constant at 0.6. This is interesting because it looks like once n is larger than 100 (could be 100,000,000!), the probability that an observation is in the bootstrap sample is still around 60%.

h.

```
store = rep(NA, 10000)
for (i in 1:10000) {
  store[i] = sum(sample(1:100, rep=TRUE) == 4) > 0
}
mean(store)
```

0.6364

Similar to what we expected, based on the chart in part **g**, four is in the bootstrap sample 63% of the time. Part **g** hypothesized that if n is greater than 100, then the probability that an observation will be in the bootstrap sample is around 60% of the time. This holds true in this case, where our observation is four and n is equal to 10,000, and likely any other that meets the "sample size" condition.

8. Section 5.4, page 198, question 5

a.

```
library(ISLR)
data(Default)
attach(Default)
set.seed(1)
glm.fit = glm(default~income+balance, family="binomial", data=Default)
```

b.

```
set.seed(1)
train = sample(nrow(Default), nrow(Default)/2)
Default.train = Default[train,]
Default.test = Default[-train,]
glm.fit = glm(default~income+balance, family="binomial", data=Default.train)
glm.probs = predict(glm.fit, Default.test, type="response")
glm.pred = ifelse(glm.probs>.5, "Yes", "No")
mean(glm.pred != Default.test$default)
```

Validation set error = 0.0286

c.

```
set.seed(5)
train = sample(nrow(Default), nrow(Default)/2)
Default.train = Default[train,]
Default.test = Default[-train,]
glm.fit = glm(default~income+balance, family="binomial", data=Default.train)
glm.probs = predict(glm.fit, Default.test, type="response")
glm.pred = ifelse(glm.probs>.5, "Yes", "No")
mean(glm.pred != Default.test$default)
```

Validation set error = 0.0246

```
set.seed(25)
train = sample(nrow(Default), nrow(Default)/2)
Default.train = Default[train,]
Default.test = Default[-train,]
glm.fit = glm(default~income+balance, family="binomial", data=Default.train)
glm.probs = predict(glm.fit, Default.test, type="response")
glm.pred = ifelse(glm.probs>.5, "Yes", "No")
mean(glm.pred != Default.test$default)
```

Validation set error = 0.0256

```
set.seed(50)
train = sample(nrow(Default), nrow(Default)/2)
Default.train = Default[train,]
Default.test = Default[-train,]
glm.fit = glm(default~income+balance, family="binomial", data=Default.train)
glm.probs = predict(glm.fit, Default.test, type="response")
glm.pred = ifelse(glm.probs>.5, "Yes", "No")
mean(glm.pred != Default.test$default)
```

Validation set error = 0.024

Even though three different train/test splits are used, the results do not vary much at all.

d.

```
set.seed(5)
train = sample(nrow(Default), nrow(Default)/2)
Default.train = Default[train,]
Default.test = Default[-train,]
glm.fit = glm(default~income+balance+student, family="binomial", data=Default.train)
glm.probs = predict(glm.fit, Default.test, type="response")
glm.pred = ifelse(glm.probs>.5, "Yes", "No")
mean(glm.pred != Default.test$default)
```

Validation set error = 0.0256

```
set.seed(25)
train = sample(nrow(Default), nrow(Default)/2)
Default.train = Default[train,]
Default.test = Default[-train,]
glm.fit = glm(default~income+balance+student, family="binomial", data=Default.train)
glm.probs = predict(glm.fit, Default.test, type="response")
glm.pred = ifelse(glm.probs>.5, "Yes", "No")
mean(glm.pred != Default.test$default)
```

Validation set error = 0.0252

```
set.seed(50)
train = sample(nrow(Default), nrow(Default)/2)
Default.train = Default[train,]
Default.test = Default[-train,]
glm.fit = glm(default~income+balance+student, family="binomial", data=Default.train)
glm.probs = predict(glm.fit, Default.test, type="response")
glm.pred = ifelse(glm.probs>.5, "Yes", "No")
mean(glm.pred != Default.test$default)
```

Validation set error = 0.024

The presence of the student dummy variable does not change much in terms of reducing the test error rate.

9. Section 5.4, page 199, question 6

a.

```
library(ISLR)
data(Default)
attach(Default)
set.seed(1)
glm.fit = glm(default~income+balance, family="binomial", data=Default)
summary(glm.fit)$coef[,1] # Same as coef(glm.fit)
```

```
> summary(glm.fit)$coef[,1]
(Intercept)      income      balance
-1.154047e+01  2.080898e-05  5.647103e-03
```

b.

```
boot.fn = function (data, index) {
  coef(glm(default~income+balance, family="binomial", data=data, subset=index))
}
```

c.

```
library(boot)
boot(Default, boot.fn, 1000)
```

```
> boot(Default, boot.fn, 1000)

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = Default, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
      original      bias      std. error
t1*  -1.154047e+01 -2.085838e-02  4.267477e-01
t2*   2.080898e-05 -1.751077e-07  4.947051e-06
t3*   5.647103e-03  1.440310e-05  2.224863e-04
```

d.

The estimated standard errors from the `glm()` and `boot.fn()` functions are very close to each other, but not exactly the same. In fact, the `glm()` function estimates are higher than the `boot.fn()` estimates! This is likely because the linear assumption causes `glm()`'s errors to be a little higher than that of bootstrap. Since bootstrap is nonparametric this assumption is not made so it is more accurate.

```
> summary(glm.fit)$coef
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.154047e+01	4.347564e-01	-26.544680	2.958355e-155
income	2.080898e-05	4.985167e-06	4.174178	2.990638e-05
balance	5.647103e-03	2.273731e-04	24.836280	3.638120e-136